

The Data Science Process

Tony Jiang

January 12, 2017

Contents

The data science process	1
The roles in a data science project	1
Stages of a data science project	1
Setting expectations	2
Loading data into R	2
Working with data from files	2
Working with relational databases	5
Exploring Data	5
Using summary statistics to spot problems	5
Spotting problems using graphics and visualization	6
Summary	22

This is study notes for “Practical Data Science with R” by Nina Zumel and John Mount, Manning 2014. -
The book: “Practical Data Science with R” by Nina Zumel and John Mount, Manning 2014 (book copyright
Manning Publications Co., all rights reserved) - The support site: [GitHub WinVector/zmPDSwR](#)

The data science process

The roles in a data science project

Project roles

A few recurring roles in a data science project

Role	Responsibilities
Project sponsor	Represents the business interest; champions the project
Client	Represents end users' interests; domain expert

Stages of a data science project

- Define the goal
- Collect and manage data
- Build the model
- Evaluate and critique model

- Present results and document
- Deploy the model to solve the problem in the real world
- Define the goal ### Defining the goal ### Data collection and management ### Modeling

Model evaluation and critique

Presentation and documentation

Model deployment and maintenance

Setting expectations

Determining lower and upper bounds on model performance

- The null model: A lower bound on performance
- The Bayes Rate: An upper bound on model performance ## Summary The data science process involves a lot of back-and-forth between the data scientist and other project stakeholders, and between the different stages of the process.

Loading data into R

Working with data from files

Work with well-structured data from files or URLs

```
uciCar <- read.csv("data/car.data.csv",sep="," ,header= TRUE)
head(uciCar) #display first 6 rows
```

```
##   buying maint doors persons lug_boot safety rating
## 1  vhigh vhigh    2        2    small    low  unacc
## 2  vhigh vhigh    2        2    small    med  unacc
## 3  vhigh vhigh    2        2    small    high unacc
## 4  vhigh vhigh    2        2     med    low  unacc
## 5  vhigh vhigh    2        2     med    med  unacc
## 6  vhigh vhigh    2        2     med    high unacc
```

```
class(uciCar)
```

```
## [1] "data.frame"
```

```
summary(uciCar)
```

```
##   buying      maint      doors      persons      lug_boot      safety
## high :432    high :432    2      :432    2      :576    big   :576    high:576
## low  :432    low  :432    3      :432    4      :576    med   :576    low  :576
## med  :432    med  :432    4      :432    more:576    small:576  med  :576
## vhigh:432    vhigh:432    5more:432
##   rating
```

```
## acc : 384
## good : 69
## unacc:1210
## vgood: 65
```

```
dim(uciCar)
```

```
## [1] 1728 7
```

Using R on less-structured data

- Transforming data in R you need “schema documentation” or “data dictionary” to decrypt troublesome data.

Read a data from German bank credit dataset

```
d <- read.table(paste('http://archive.ics.uci.edu/ml/',
  'machine-learning-databases/statlog/german/german.data', sep=''),
  stringsAsFactors=F, header=F)
print(d[1:3,])
```

```
##      V1 V2  V3  V4   V5  V6  V7 V8  V9  V10 V11  V12 V13  V14  V15 V16  V17
## 1 A11  6 A34 A43 1169 A65 A75  4 A93 A101  4 A121  67 A143 A152  2 A173
## 2 A12 48 A32 A43 5951 A61 A73  2 A92 A101  2 A121  22 A143 A152  1 A173
## 3 A14 12 A34 A46 2096 A61 A74  2 A93 A101  3 A121  49 A143 A152  1 A172
##      V18  V19  V20 V21
## 1      1 A192 A201  1
## 2      1 A191 A201  2
## 3      2 A191 A201  1
```

Change the column names to something meaningful

```
colnames(d) <- c('Status.of.existing.checking.account',
  'Duration.in.month', 'Credit.history', 'Purpose',
  'Credit.amount', 'Savings account/bonds',
  'Present.employment.since',
  'Installment.rate.in.percentage.of.disposable.income',
  'Personal.status.and.sex', 'Other.debtors/guarantors',
  'Present.residence.since', 'Property', 'Age.in.years',
  'Other.installment.plans', 'Housing',
  'Number.of.existing.credits.at.this.bank', 'Job',
  'Number.of.people.being.liable.to.provide.maintenance.for',
  'Telephone', 'foreign.worker', 'Good.Loan')
d$Good.Loan <- as.factor(ifelse(d$Good.Loan==1, 'GoodLoan', 'BadLoan'))
print(d[1:3,])
```

```
##      Status.of.existing.checking.account Duration.in.month Credit.history
## 1                                     A11                      6          A34
## 2                                     A12                      48          A32
## 3                                     A14                      12          A34
##      Purpose Credit.amount Savings account/bonds Present.employment.since
```

```
## 1      A43      1169      A65      A75
## 2      A43      5951      A61      A73
## 3      A46      2096      A61      A74
##      Installment.rate.in.percentage.of.disposable.income
## 1      4
## 2      2
## 3      2
##      Personal.status.and.sex Other.debtors/guarantors Present.residence.since
## 1      A93      A101      4
## 2      A92      A101      2
## 3      A93      A101      3
##      Property Age.in.years Other.installment.plans Housing
## 1      A121      67      A143      A152
## 2      A121      22      A143      A152
## 3      A121      49      A143      A152
##      Number.of.existing.credits.at.this.bank Job
## 1      2 A173
## 2      1 A173
## 3      1 A172
##      Number.of.people.being.liable.to.provide.maintenance.for Telephone
## 1      1      A192
## 2      1      A191
## 3      2      A191
##      foreign.worker Good.Loan
## 1      A201      GoodLoan
## 2      A201      BadLoan
## 3      A201      GoodLoan
```

Building a map to interpret loan use codes

```
mapping <- list(
  'A40'='car (new)',
  'A41'='car (used)',
  'A42'='furniture/equipment',
  'A43'='radio/television',
  'A44'='domestic appliances' # note that other codes are not defiend here.
)
```

Transform the data

```
for(i in 1:(dim(d))[2]) { # Note: 1
  if(class(d[,i])=='character') {
    d[,i] <- as.factor(as.character(mapping[d[,i]])) # Note: 2
  }
}
table(d$Purpose,d$Good.Loan)
```

```
##
##      BadLoan GoodLoan
## car (new)      89      145
## car (used)     17      86
## domestic appliances      4      8
## furniture/equipment     58     123
```

```
##      NULL                70      120
##    radio/television      62      218
```

Working with relational databases

The right way to work with data found in databases is to connect R directly to the database.

A production-size example

- United States Census 2011 national PLUMS American Community Survey data (www.census.gov/acs/www/data_documentation)
- Millions of rows
- a few gigabytes when zipped
- This size is the sweet spot for relational database or SQL database. we are not forced to move into a MapReduce or database cluster to do our work.

Curating the data Staging the data into a database - H2 - SQL Screwdriver Squirrel SQL

Exploring Data

Resist the temptation to dive into the modeling step without looking at the dataset first.

Using summary statistics to spot problems

```
custdata<- read.csv('data/custdata.tsv',header=T,sep='\t')
summary(custdata)
```

```
##      custid      sex  is.employed      income
##  Min.   : 2068  F:440  Mode :logical  Min.   : -8700
##  1st Qu.: 345667 M:560  FALSE:73   1st Qu.: 14600
##  Median : 693403      TRUE :599   Median : 35000
##  Mean   : 698500      NA's :328   Mean   : 53505
##  3rd Qu.:1044606                      3rd Qu.: 67000
##  Max.   :1414286                      Max.   :615000
##
##      marital.stat health.ins
##  Divorced/Separated:155  Mode :logical
##  Married           :516  FALSE:159
##  Never Married     :233  TRUE :841
##  Widowed           : 96  NA's :0
##
##
##      housing.type recent.move      num.vehicles
##  Homeowner free and clear :157  Mode :logical  Min.   :0.000
##  Homeowner with mortgage/loan:412  FALSE:820   1st Qu.:1.000
##  Occupied with no rent    : 11  TRUE :124   Median :2.000
##  Rented                   :364  NA's :56    Mean   :1.916
```

```
## NA's : 56 3rd Qu.:2.000
## Max. :6.000
## NA's :56
## age state.of.res
## Min. : 0.0 California :100
## 1st Qu.: 38.0 New York : 71
## Median : 50.0 Pennsylvania: 70
## Mean : 51.7 Texas : 56
## 3rd Qu.: 64.0 Michigan : 52
## Max. :146.7 Ohio : 51
## (Other) :600
```

Typical problems revealed by data summaries

- Missing values
- Invalid values and outliers
- Data range: relative because of units. (babies age in weeks vs years)
- Units

Spotting problems using graphics and visualization

We cannot expect a small number of numerical values to consistently convey the wealth of information that

The use of graphics to examine data is called visualization.

Visually checking distributions for a single variable

what is the peak value? How many peaks ? How normal is the data? How much does the data vary?

- Histogram

```
library(ggplot2)
ggplot(custdata)+geom_histogram(aes(x=age),binwidth=5,fill="gray")
```

- Density plots

```
library(scales)
ggplot(custdata)+geom_density(aes(x=income))+scale_x_continuous(labels=dollar)
```






Figure 1: A histogram tells you where your data is concentrated



```
ggplot(custdata)+geom_density(aes(x=income))+scale_x_log10(breaks=c(100,1000,10000,100000),labels=dollars)
```

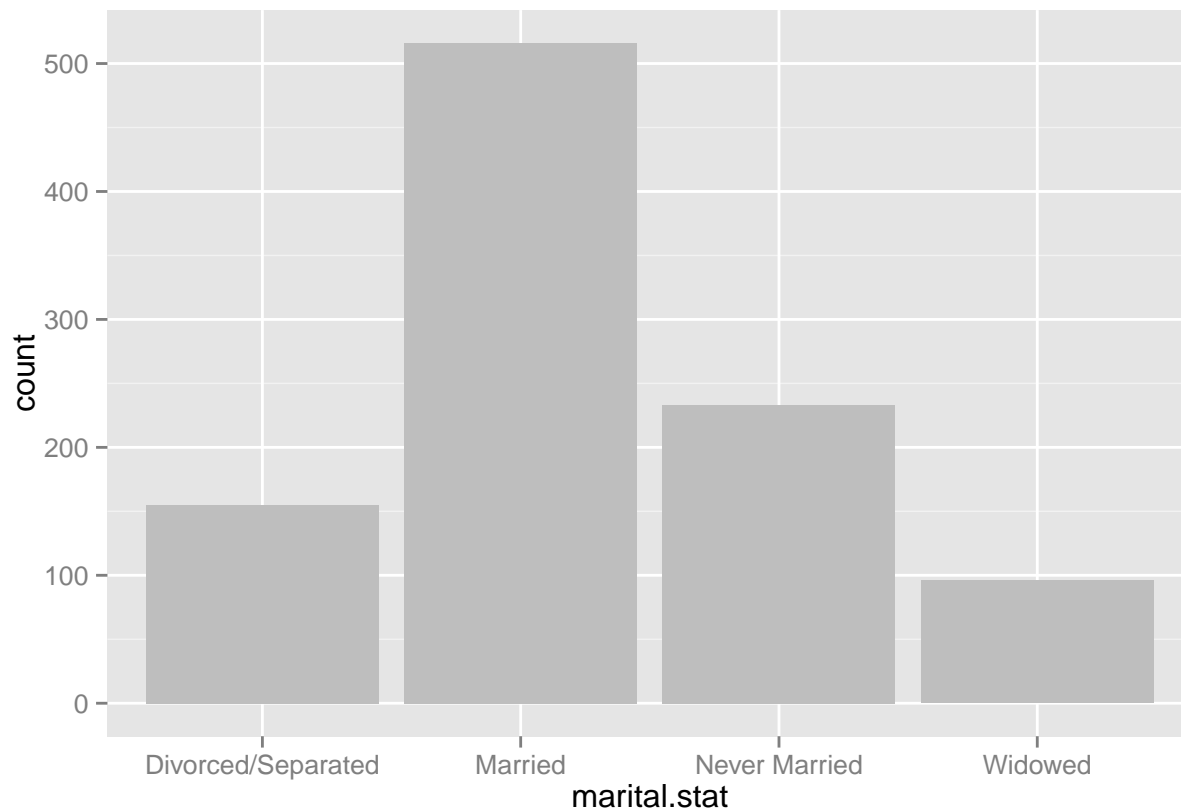
```
## Warning in scale$trans$trans(x): NaNs produced
```

```
## Warning: Removed 79 rows containing non-finite values (stat_density).
```

- Bar charts


```
ggplot(custdata)+geom_bar(aes(x=marital.stat),fill="gray")
```



A


horizontal bar chart can be easier to read when there are several categories with long names.

```
ggplot(custdata)+geom_bar(aes(x=state.of.res),fill="gray")+coord_flip()+theme(axis.text.y=element_text(
```



Cleveland recommends that the data in a bar chart be sorted, to more efficiently extract insight from the data.


```
statesums<-table(custdata$state.of.res)
statef<-as.data.frame(statesums)
colnames(statef)<-c("state.of.res","count")
statef<-transform(statef,state.of.res=reorder(state.of.res,count))
ggplot(statef)+geom_bar(aes(x=state.of.res,y=count),stat="identity",fill="gray")+coord_flip()+theme(axis
```



Visually checking relationships between two variables

- Line Plots

```
x<-runif(100)
y<-x^2+0.2*x
ggplot(data.frame(x=x,y=y),aes(x=x,y=y))+geom_line()+geom_point()
```




Scatter Plots and smoothing curves

```
custdata2<-subset(custdata,(custdata$age>0 & custdata$age<100 & custdata$income>0))  
cor(custdata2$age,custdata2$income)
```

```
## [1] -0.02240845
```

```
ggplot(custdata2,aes(x=age,y=income))+geom_point()+ylim(0,200000)
```

```
## Warning: Removed 32 rows containing missing values (geom_point).
```




Linear fit

```
ggplot(custdata2,aes(x=age,y=income))+geom_point()+geom_smooth(method="lm")+ylim(0,200000)
```

```
## Warning: Removed 32 rows containing missing values (stat_smooth).
```

```
## Warning: Removed 32 rows containing missing values (geom_point).
```




Smoothing Curve

```
ggplot(custdata2,aes(x=age,y=income))+geom_point()+geom_smooth()+ylim(0,200000)
```

```
## geom_smooth: method="auto" and size of largest group is <1000, so using loess. Use 'method = x' to c
```

```
## Warning: Removed 32 rows containing missing values (stat_smooth).
```


```
## Warning: Removed 32 rows containing missing values (geom_point).
```



Distribution of customers with health insurance, as a function of age

```
ggplot(custdata2,aes(x=age,y=as.numeric(health.ins)))+geom_point(position=position_jitter(w=0.05,h=0.05))
```

```
## geom_smooth: method="auto" and size of largest group is <1000, so using loess. Use 'method = x' to cl
```




- Hexbin Plots

```
library(hexbin)
ggplot(custdata2, aes(x=age, y=income)) + geom_hex(binwidth=c(5, 10000)) + geom_smooth(color="white", se=F) + ylim(0, 100000)

## Warning: Removed 32 rows containing missing values (stat_hexbin).

## geom_smooth: method="auto" and size of largest group is <1000, so using loess. Use 'method = x' to choose 'lm' or 'glm'


## Warning: Removed 32 rows containing missing values (stat_smooth).
```


- Bar charts for two categorical variables

Relationship between marital status and the probability of health insurance coverage.


```
# bar chart  
ggplot(custdata)+geom_bar(aes(x=marital.stat,fill=health.ins))
```



```
ggplot(custdata)+geom_bar(aes(x=marital.stat,fill=health.ins),position="dodge")
```

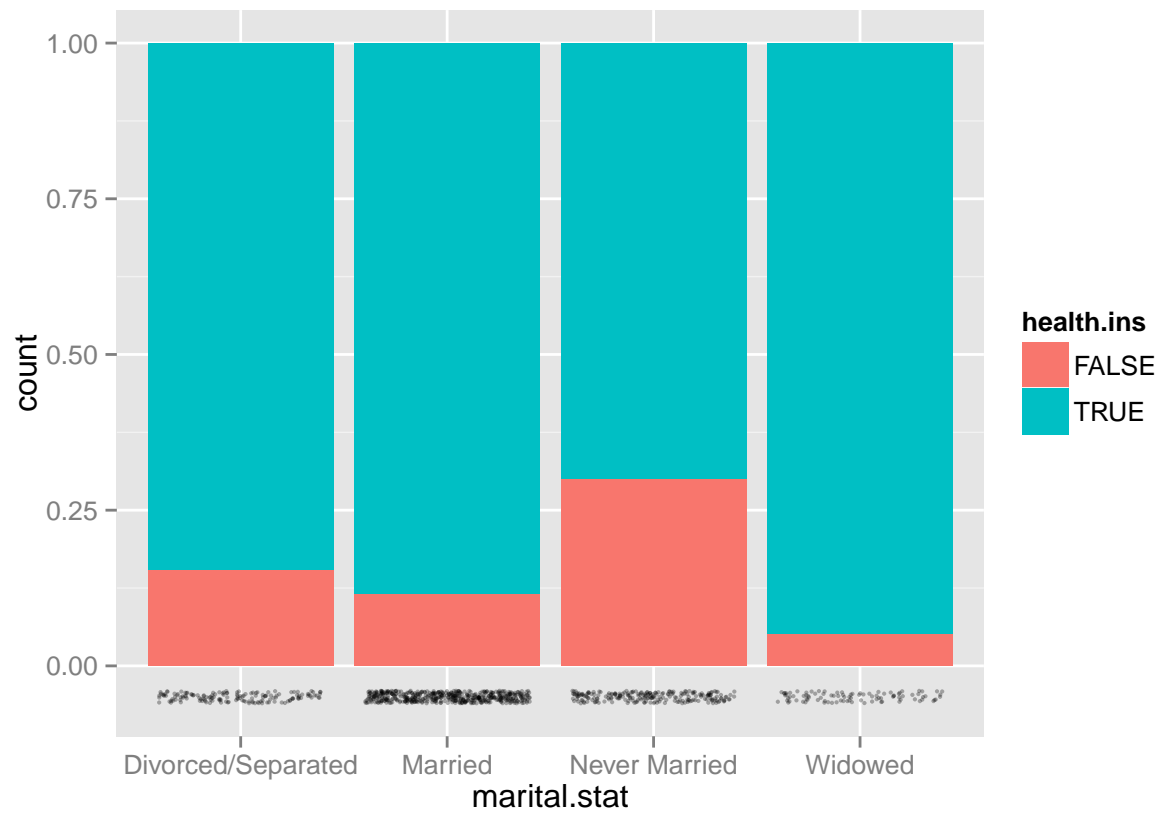


```
ggplot(custdata)+geom_bar(aes(x=marital.stat,fill=health.ins),position="fill")
```




add a rug

```
ggplot(custdata,aes(x=marital.stat))+geom_bar(aes(fill=health.ins),position="fill")+geom_point(aes(y=-0
```

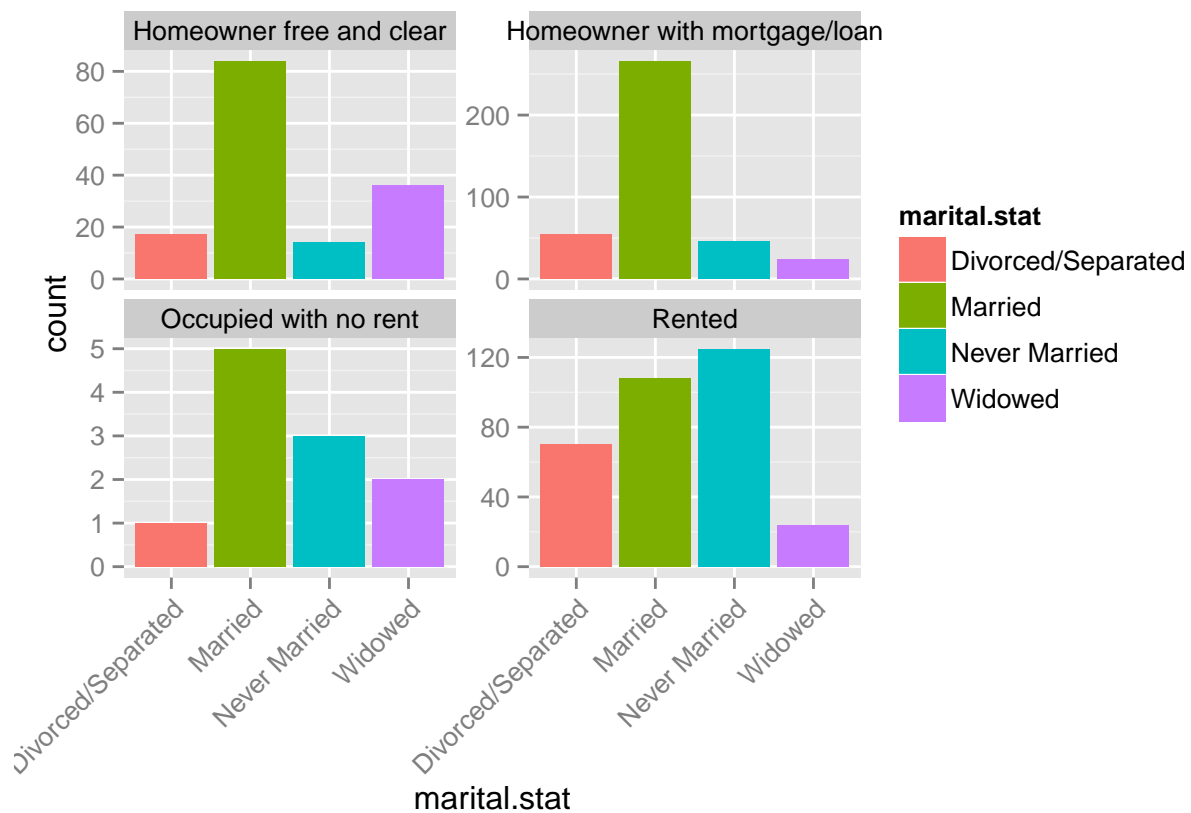


Facetd

```
ggplot(custdata2)+geom_bar(aes(x=housing.type,fill=marital.stat),position="dodge")+theme(axis.text.x=el
```



```
ggplot(subset(custdata2, housing.type != "NA")) + geom_bar(aes(x = marital.stat, fill = marital.stat), position = "dodge")
```



Summary

- Take the time to examine your data before diving into the modeling
- The `summary` command helps you spot issues with data range, units, data type, and missing or invalid values
- Visualization additionally gives you a sense of data distribution and relationships among variables
- Visualization is an iterative process and helps answer questions about the data. Time spent here is time not wasted during the modeling process.